

---

# T1: Secure Programming for Embedded Systems

NorthSec 2019

---

# Developing for Embedded Systems

# What is an “Embedded System”?

---



**Embedded**



**Not embedded**

# What is an “Embedded System”?

---

Definition of “embedded system” is arbitrary.

What is meant here:

- Small 16-bit or 32-bit CPU (e.g. ARM Cortex M0+)
- RAM: 64 kB or less
- ROM: 256 kB or less
- Some network connectivity
- No operating system (“bare metal”)
- Strong constraints on size / power / thermal dissipation

(CPU is feeble but this does not matter much.)

---

# Constraints: Consequences on Security

---

## No memory management unit (MMU)

- All RAM is accessible read/write (and exec in some architectures)
- ROM (Flash) is all readable
- No sandbox / isolation
- No trapping of NULL pointer dereference
- No ASLR
- No guard page for stack overflows
  - Recursive algorithms must be banned

# Constraints: Consequences on Security

---

## No room for multiple or large stacks

- Multiple concurrent processes must run
- ... but without locking the system
- A typical C stack needs at least 1-2 kB, more realistically 4 kB
- C tends to increase stack usage

# Constraints: Consequences on Security

---

```
static
void battery_status_timeout_handler(void *p_context) {
    char msg[256];

    gfx_fillRect(0, 8, 128, 56, SSD1306_BLACK);
    gfx_setCursor(0, 12);
    gfxSetTextBackgroundColor(SSD1306_WHITE, SSD1306_BLACK);

    snprintf(msg, sizeof(msg),
        "Battery status:\n"
        " Voltage: %04d mV\n"
        " Charging: %s\n"
        " USB plugged: %s\n",
        battery_get_voltage(),
        battery_is_charging() ? "Yes" : "No",
        battery_is_usb_plugged() ? "Yes" : "No");

    gfx_puts(msg);
    gfx_update();
}
```

# Constraints: Consequences on Security

---

0:	e92d 41f0	stmdb sp!, {r4, r5, r6, r7, r8, lr}	<b>24 bytes</b>
4:	b0c2	sub sp, #264 ; 0x108	<b>264 bytes</b>
6:	2400	movs r4, #0	
8:	2338	movs r3, #56 ; 0x38	
a:	2280	movs r2, #128 ; 0x80	
c:	4620	mov r0, r4	
e:	af02	add r7, sp, #8	
10:	9400	str r4, [sp, #0]	
12:	2108	movs r1, #8	
14:	f7ff fffe	bl 0 <gfx_fillRect>	
18:	4620	mov r0, r4	
1a:	210c	movs r1, #12	
1c:	f7ff fffe	bl 0 <gfx_setCursor>	
20:	4621	mov r1, r4	
22:	2001	movs r0, #1	
24:	f7ff fffe	bl 0 <gfxSetTextBackgroundColor>	
( . . . )			

# Languages for embedded development

---

## C

- Works everywhere
- “Portable assembly” but with a few hidden automatic costs
- Not *memory-safe*:
  - No check on array accesses
  - Manual allocation / deallocation → double-free, use-after-free, leaks...
  - Type punning
- “Undefined Behavior”
- Often required at some level (e.g. SDK offers only a C API)
  - It’s a C world

# Languages for embedded development

## Java ME

- GC, strong types,...
- Large RAM / ROM requirements
- Only ARM
- Needs an OS

**Q:** What are the system requirements for Oracle Java ME Embedded 8?

**A:** The high-level system requirements are as follows:

- System based on ARM architecture SOCs
- Memory footprint as low as 128 KB RAM and 1 MB ROM (see note)
- Very simple embedded kernel, or a more capable embedded OS/RTOS
- At least one type of network connection (wired or wireless)

**Note:** Footprint based on MEEP 8 Minimal Profile Set, optimized for single-function devices. Actual footprint will vary based on target device and use case.

# Languages for embedded development

---

## Go

- Only with TinyGo: <https://tinygo.org/>
- Limited language / runtime support:
  - “*support for goroutines and channels is weak*”
  - Maps can only have up to 8 (eight!) entries
  - GC: only for ARM, other platforms “*will just allocate memory without ever freeing it*” (but GC is required for proper string management)

# Languages for embedded development

---

## Rust Embedded: <https://www.rust-lang.org/what/embedded>

- Inherits all the memory-safety features of Rust
- Heap is optional
  - But without the heap, everything is allocated on the stack
- Supports ARM Cortex-M and Cortex-R, RISC-V, and MSP430 (experimental)
  - But not AVR or Xtensa or other architectures that LLVM does not support
- Typically more stack-hungry than C
- Lots of automatic magic

# Languages for embedded development

---

## Forth

- Many incompatible implementations
  - It's more a concept than a single defined language (though there is an ANSI standard)
  - You are supposed to “write your own Forth”
- Very compact, with low RAM usage
- Even less safe than C, and extremely non-portable

# Languages for embedded development

---

## Summary:


- No perfect language
- Adaptations from “larger languages” don’t solve the inherent issues, especially the cost of stacks for concurrent processing
- Often needs to interoperate with C
- Generic portability requires compiling *to* C
- Security is better addressed with a *non-magic* language

---

# Success Story: BearSSL and T0

# BearSSL

---



SSL/TLS library optimized for embedded systems


- Full-featured with uncompromising security (e.g. constant-time code)
- Portable, no dependency on any specific runtime, OS or compiler
- State-machine API
- No dynamic memory allocation whatsoever
- Can run in limited ROM and RAM (about 21 kB ROM and 25 kB RAM)
  - Can use less RAM, but requires support of small records by the peer

**Problem:** TLS handshake messages, and X.509 certificates, are complex, nested structures that can be large.

- X.509 certificate chain can be up to 16 MB
  - Realistically, 2 to 10 kB; sometimes larger (OpenSSL's default max is 100 kB)
- Data can be fragmented over different records
- Cannot buffer a complete message or certificate
  - Must perform streamed processing
  - Processing must be interruptible and restartable

**Idea:** run the decoding process as a *coroutine*

# BearSSL



- BearSSL is computational only (application handles low-level I/O)
- Handshake parser and X.509 validation run as two coroutines
  - Each has its own state (stacks, variables)
  - Parsing proceeds when data becomes available, by chunks

# T0

---

**T0** is a Forth-like language used to implement the handshake parser and the X.509 validation engine.

- Compiled to *threaded code*
- Uses two custom stacks (data & system stack) of limited size (128 bytes each)
- Runs in a flat, small interpreter loop that can be stopped and restarted at will
- Instructions are a single byte each (*token threading*)
- Compiler is written in C# and performs some static analysis (maximum stack usage)

# Threaded Code

---

```
\ Read one byte, enforcing current read limit.  
: read8 ( lim -- lim x )  
    dup ifnot ERR_X509_INNER_TRUNC fail then  
        1- read8-nc ;  
  
\ Read a 16-bit value, big-endian encoding.  
: read16be ( lim -- lim x )  
    read8 8 << swap read8 rot + ;  
  
\ Read a 16-bit value, little-endian encoding.  
: read16le ( lim -- lim x )  
    read8 swap read8 8 << rot + ;
```

Executable code is (mostly) a sequence of function calls.

---


# Indirect Threaded Code

```
: read8 ( lim -- lim x )
    dup ifnot 36 fail then 1- read8-nc ;
: read16be ( lim -- lim x )
    read8 8 << swap read8 rot + ;
```

```
for (;;) {
    w = *ip++;
    (*w)();
}
```

*Virtual CPU*

```
CALL read8
CONST 8
CALL <<
CALL swap
CALL read8
CALL rot
CALL +
RET
```



# Indirect Threaded Code

---

- Each function is a memory structure whose first field (CFA) is a pointer to native code.
  - For *primitive* functions, there is only that pointer.
  - *Interpreted* functions use the generic entry code ({ENTER}); CFA is followed by the function code as a sequence of pointers to function structures.
  - Some primitive functions extract arguments located in the calling code (e.g. local jumps).
  - Execution proceeds with a virtual CPU loop and two stacks:
    - *Data stack*: for function arguments and returned values
    - *Return stack*: for return addresses and local variables
- **Stack usage is explicit**

# Token Threaded Code

---

- Each pointer to a function structure is replaced with a token (index in a table of pointers).
- One extra indirection per instruction.
- Most/all instructions fit on one byte.
- Primitive function code can be integrated inside the virtual CPU loop.

# T0 Compilation

---

```
$ ./T0Comp.exe -o src/x509/x509_minimal -r br_x509_minimal src/x509/asn1.t0 src/x509/x509_minimal.t0  
[src/x509/asn1.t0]  
[src/x509/x509_minimal.t0]  
main: ds=17 rs=25  
code length: 2836 byte(s)  
data length: 299 byte(s)  
total words: 203 (interpreted: 142)
```

- Compiler reads and interprets T0 code
  - *Immediate functions* are executed on-the-fly (metaprogramming)
- C source code is produced with tokens, primitives and virtual CPU
- X.509 validator compiled size (ARM Cortex M4):

```
$ size x509_minimal.o  
text      data      bss      dec      hex filename  
6259        0        0    6259    1873 x509_minimal.o
```

# T0 Advantages

---

- Code can run as a coroutine with very small state (168 bytes for the two stacks)
- No dynamic memory allocation; streamed processing
- Guaranteed maximum stack usage
- Compiler verifies “types” (stack depth at all points)
- Small code footprint
- *No magic*
  
- ... but not completely *memory-safe*

---

T1

Evolution of T0 with extra features:

- Memory-safe
- Optional dynamic memory allocation (controlled) with GC
- Rich type system (including generics)
- OOP support
- Namespaces and modules

# Memory Safety

---

*Memory safety* is a set of memory-related features:

- No uncontrolled type punning
- Array accesses outside of bounds are prevented
- No use-after-free or double-free
- Guaranteed stack usage (no overflow)
- Guaranteed maximum heap usage
- All allocated memory is released (no leak)
- Concurrent writing is controlled or prevented
- Etc...

# Memory Safety in T1

---

Runtime checks:

- Array bounds on access
- Automatic memory management (garbage collector)

Compile-time checks:

- Maximum stack sizes
- Escape analysis (for stack-allocated objects)
- All method lookups are solvable
- No memory is interpreted with the wrong type
- No write access to static constant objects

# OOP

---

```
class A {  
    void foo(A a) {  
        System.out.println("foo AA");  
    }  
    void foo(B b) {  
        System.out.println("foo AB");  
    }  
}  
class B extends A {  
    void foo(A a) {  
        System.out.println("foo BA");  
    }  
    void foo(B b) {  
        System.out.println("foo BB");  
    }  
}  
class C {  
    public static void main(String[] args) {  
        A x = new B();  
        A y = new B();  
        x.foo(y);  
    }  
}
```

## Java code:

- Method call has a special first parameter (*object on which the method is called*)
- Method lookup uses the dynamic (runtime) type of the first parameter
- For other parameters, the static (compile-time) type is used

→ This program prints:

**foo BA**

# OOP

---

```
struct A
end

struct B <sub> A
end

: foo (A A)
  "foo AA" println ;
: foo (A B)
  "foo AB" println ;
: foo (B A)
  "foo BA" println ;
: foo (B B)
  "foo BB" println ;

: main ()
  B new B new ->{ x y }
  x y foo ;
```

## T1 code:

- No special parameter
- Method lookup uses the dynamic types of all parameters
- No explicit static type analysis

→ This program prints:

**foo BB**

# Types

---

- Each value is a *pointer*
  - Plain integers, Booleans... are also “pointers”
  - No “value type”
- Every access to an object field is through an *accessor* (dedicated method)
  - Accessors locate the field unambiguously
- Basic types:
  - Booleans: bool
  - Plain integers: int
  - Modular integers: u8 u16 u32 u64 i8 i16 i32 i64

# NULL

---

There is no null pointer value.

- Reading from an uninitialized object field triggers a runtime error
- Some object fields (basic types) are initialized at zero
- Possible reads from uninitialized local variables are detected at compilation

# Plain Integers and Overflows

---

Strategies when integer operations overflow the representable range:

- Use modular arithmetic (C#, Java, Go)
- Report an error (Ada)
- Do one or the other, depending on external circumstances (Rust)
- Transparently upgrade to big integers (Python, Scheme)
- Use floating point (JavaScript)
- Anything goes (C, C++)

T1 uses the Ada way for “plain integers” (`int`) and modular arithmetic for exact-width integers (`u16`, `i32...`)


---

# Whole Program Analysis

```
: triple (object)
    dup dup + + ;

: main ()
    4i32 triple println
    "foo" triple println ;
```

- Compute the complete call tree with possible stack contents.
- Each call of a given function is a different node.



# Whole Program Analysis

---

- Complete flow analysis from entry point:
  - For each function call, only cares about which types can actually be present on the stack.
  - Types for function definition are for call routing, not type restriction.
  - No syntax to express *potential* parameter types.
  - Return types are computed.
  - Dead opcodes and unreachable functions are detected.
- Multiple nodes for each function (one per call site):
  - All functions are *generic*.
  - Recursion would lead to an infinite tree (disallowed).
- Includes escape analysis and detection of writes to constant instances.

# Current Status

---

**Web site:** <https://t1lang.github.io/>

## Done:

- Specification + rationale
- Bootstrap interpreter/compiler:
  - Interpreter
  - Whole program analysis
  - Code generator (partial)

## TODO:

- Finish bootstrap compiler
- Standard library (at least lists and sorted maps)
- Rewrite T1 compiler in T1